

Handwritten Character Classification Using Convolutional Neural Net

By: Chris Corona and
Nolan Walker

Motivation

Computer vision

- Video / image classification
- Object recognition
- Recommendation systems
- Natural language processing

Convolutional Neural Network

A neural net that is especially good with images because it has a built-in filtering mechanism.

- Assign importance to various aspects in an image
- Captures spatial and temporal dependencies
- Learn high-level features

Convolutional Layer

Apply kernel to obtain convolved feature

- Reduce image into form that is easier to process, without losing features

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image



4		

Convolved Feature

Padding and Stride

Change in dimensionality
from choice of padding and
stride

- Padding for same / increased dimensionality
- No padding for reduced dimensionality
- Stride for more reduction




Pooling Layer

Way to reduce a layer's complexity, downsample the number of features conveyed

- Max
- Average
- Global: at the end

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0


3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1



Activation Function

Introduce non-linearity to the model

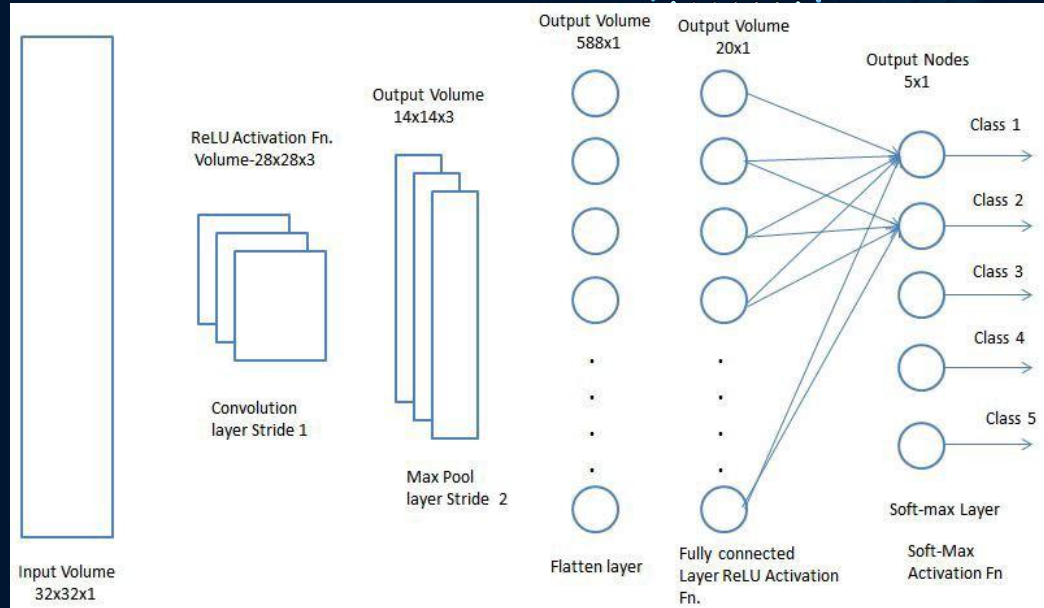
- Sigmoid
- Tanh
- ReLU
- LeakyReLU



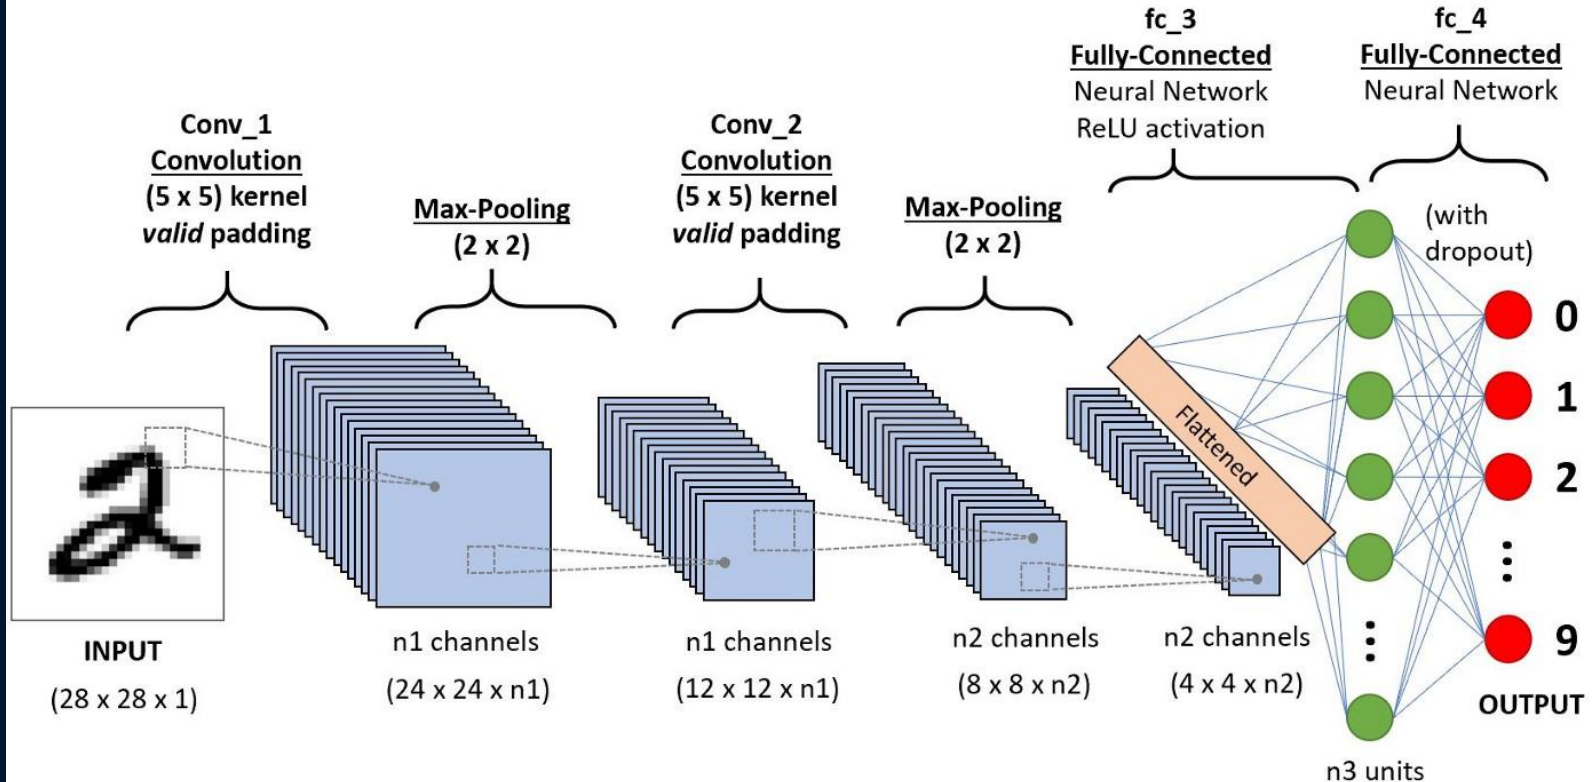
Classification Layer

Flatten final output from convolutional layers and send to regular neural net

- Learn combinations of high-level features
- Possibly non-linear functions
- Softmax layer



CNN Sequence




EMNIST Dataset

The **Extended** MNIST (EMNIST) dataset is the successor to the popular MNIST dataset

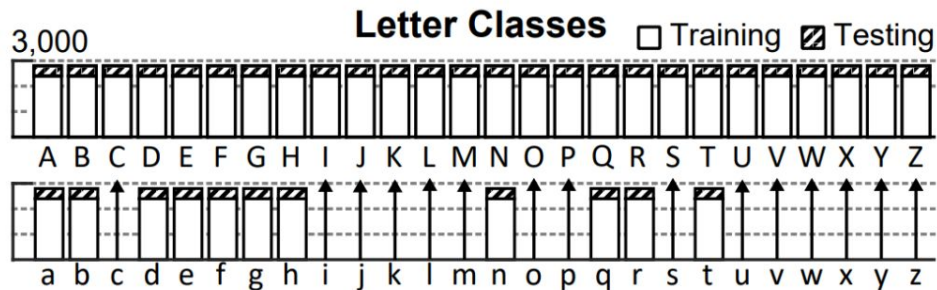
- 131,600 images, 28 x 28 pixels in grayscale
 - 112,800 train / 18,800 test
- Pre-processed
 - Size-normalized, centered, and fixed-size image
- 47 balanced classes
 - Digits: 0-9
 - Uppercase letters: A-Z
 - Lowercase letters: a, b, d, e, f, g, h, n, q, r, t
- More challenging than MNIST

EMNIST Dataset

EMNIST Balanced Dataset



47 Classes, 131,600 Samples




EMNIST Dataset

1 4 X E h 1 W Q k 9 9 a D H Q P A V a O R H f B h
1 5 H G h b 1 C D 8 a h e F V U P Q 9 b 3 4 5 Q m
h a e 5 6 a K G # 1 P 5 6 9 m 4 u a b H 8 1 h e m
n w v o o n + 2 5 X B 6 0 v 2 0 u e H m + y ~ c a
J 3 E F Z W S f W w 2 b 4 J T j B P Z D a 2 v o B
Y f Q J d e f a b 3 8 G J T S j X Y Y m Q b Y 7 d
o b d f z o 5 o n f 4 1 b 5 0 7 w a o m r B n B P
+ 9 h A k z L E e n i m K F O B H X v a v f G W A
h 3 # N Q t s q w C Z F d a w 4 6 s n a 8 4 f G A
4 7 R E G + e C Y K P e f U R 9 s H Z b o T S n X
w 2 H R e e 9 f r o Q f y 9 B P J n M i t p r s s
4 1 J o Z G r z y 2 x G d 4 6 s B 3 8 A + t i k 7
w F H B f o u 1 E Y C 9 H o u N o b y 4 n 2 3 H 2
h 6 b P N 9 u 5 3 n 9 z F y w 5 A d Y B 4 P a i 4
9 f Q 9 C n 3 C 6 d n 5 6 T T e 1 6 D O P 3 f 3 5

Baseline Results

In 2017, Gregory Cohen created an instructive baseline with his paper *EMNIST: An extension of MNIST to handwritten letters*

- OPIUM classifier
 - Accuracy 78.02%
 - 10,000 hidden nodes
- Linear classifier (not shown)
 - Accuracy 50.93%



Our Results

Using a very basic convolutional neural net, we were able to exceed the Cohen et al results from 2017

- CNN classifier
 - Accuracy 87.22%
 - 3 convolutional layers
 - ReLu activation
 - 2 max pooling layers

Our Results

```
layers = [  
    imageInputLayer([28,28,1])  
  
    convolution2dLayer(3,8,'Padding','same')  
    batchNormalizationLayer  
    reluLayer  
  
    maxPooling2dLayer(2,"Stride",2)  
  
    batchNormalizationLayer  
    convolution2dLayer(3,16,'Padding','same')  
    reluLayer  
  
    maxPooling2dLayer(2,"Stride",2)  
  
    convolution2dLayer(3,16,'Padding','same')  
    batchNormalizationLayer  
    reluLayer  
  
    fullyConnectedLayer(47)  
    softmaxLayer  
    classificationLayer  
];
```

Results

Validation accuracy:	87.22%
Training finished:	Max epochs completed

Training Time

Start time:	26-Apr-2022 15:36:53
Elapsed time:	85 min 11 sec

Training Cycle

Epoch:	20 of 20
Iteration:	17620 of 17620
Iterations per epoch:	881
Maximum iterations:	17620

Validation

Frequency:	50 iterations
------------	---------------

Other Information

Hardware resource:	Single CPU
Learning rate schedule:	Constant
Learning rate:	0.01



Demonstration

Need volunteers to suggest characters

Characters will be drawn in MS Paint

Use the trained net to make live predictions

Can we get 8 out of 9 correct? ($\approx 88\%$)

References

Cohen, G., Afshar, S., Tapson, J., Schaik, A. (2017). EMNIST: an extension of MNIST to handwritten letters. *2017 International Joint Conference on Neural Networks*. 2921-2926.

<https://www.nist.gov/itl/products-and-services/emnist-dataset>

<https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53>